YOULearn.com(Your Online University)

Cloud Data Architecture

INDEX

1. INTRODUCTION:

Online education platforms process large amounts of data. ***YOULearn.com*** manages user profiles, course content, payments, and engagement metrics. A well-structured cloud data system ensures efficient storage, retrieval, and analysis.

This case study outlines a cloud-based architecture that integrates multiple data sources, processes them, and generates reports and insights. The goal is to improve platform performance, enhance security, and provide useful data for learners, instructors, and administrators.

2. Objectives:

Build a scalable cloud system that integrates multiple data sources. Provide real-time insights for students, instructors, and administrators. Ensure System reliability, security, and performance through monitoring.

3. Design Phase of Cloud Architecture

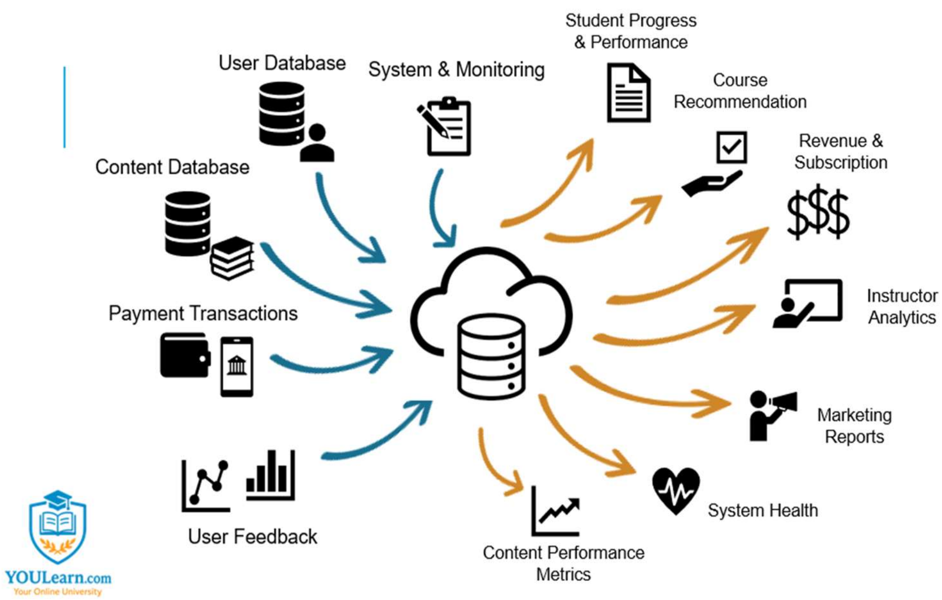

DATABASE INPUT

User Database

- **Description:** Stores student and instructor profiles, enrollments, and progress.
- **Location:** AWS RDS (PostgreSQL/MySQL), accessible via API.
- **Type:** Structured, SQL tables.

Course Content Storage

- **Description:** Stores videos, documents, quizzes, and interactive content.
- **Location:** AWS S3/Azure Blob, accessible via API.
- **Type:** Unstructured, multimedia files.

Payment Transactions

- **Description:** Stores subscription payments, refunds, and invoices.

- **Location:** AWS RDS, integrated with Stripe/PayPal APIs.
- **Type:** Structured, SQL tables.

User Engagement Data

- **Description:** Tracks course completion, time spent, quiz scores, and activity.
- **Location:** Google BigQuery/AWS Redshift.
- **Type:** Structured, SQL tables

System Logs

- **Description:** Captures system events, errors, and performance metrics.
- **Location:** AWS CloudWatch, ELK Stack.
- **Type:** Semi-structured, log files.

DATABASE OUTPUT

Student Reports & Certifications

- **Delivery:** Auto-generated reports and dashboards.
- **Used by:** Students, instructors, administrators.
- **Example:** A course completion certificate.

Course Recommendations

- **Delivery:** AI-powered recommendation engine.
- **Used by:** Students.
- **Example:** "You might like 'Python for Beginners.'"

Revenue Reports

- **Delivery:** Financial dashboards and reports.
- **Used by:** Business and finance teams.
- **Example:** Monthly subscription revenue breakdown.

Content Performance Metrics

- **Delivery:** Reports and analytics dashboards.
- **Used by:** Instructors, content teams.
- **Example:** Dropout rates for each lecture.

Instructor Dashboards

- **Delivery:** Web-based analytics dashboard.
- **Used by:** Instructors.

- **Example:** Course completion rates and student engagement.

Marketing Insights

- **Delivery:** AI-driven reports.
- **Used by:** Marketing teams.
- **Example:** User retention trends.

System Health Reports

- **Delivery:** Real-time monitoring dashboards.
- **Used by:** IT and DevOps teams.
- **Example:** Server uptime and error logs.

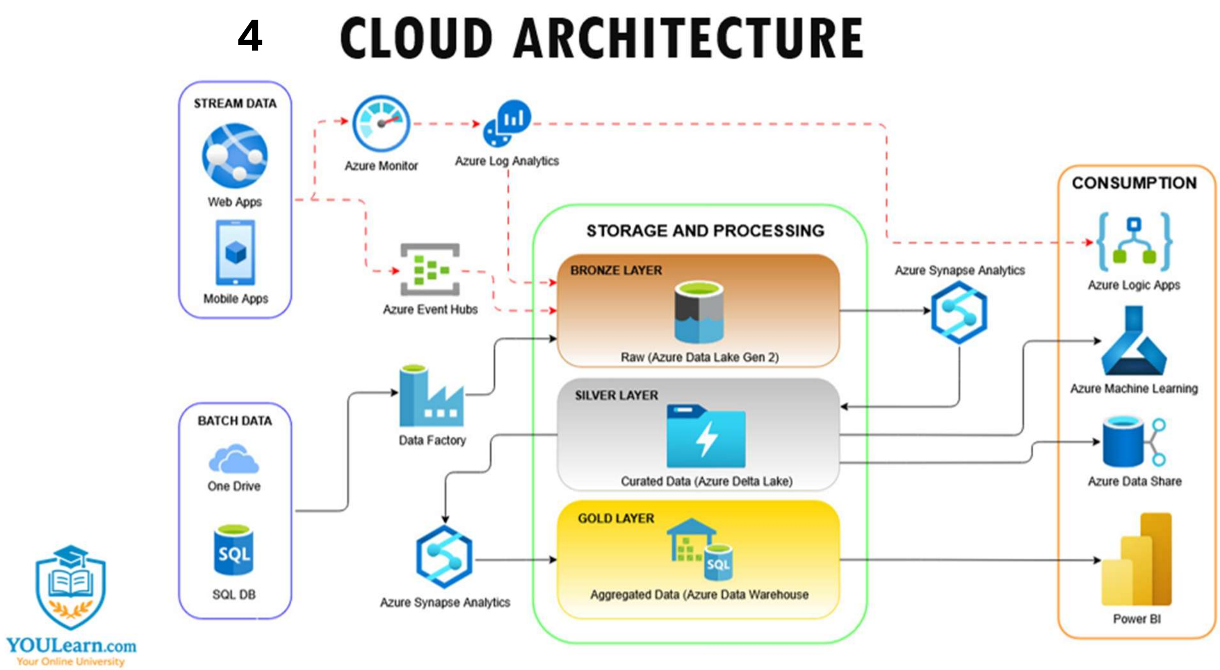

Simple Flow Summary:

Data → Monitor/Move → Store (Raw → Clean → Final) → Use the data for apps, AI, sharing, and reports

Step 1: Data Sources

Where your university's data comes from:

- **Stream Data (LIVE activities):**
    - **Students using your website and mobile app.**
    - **Example: A student signs up for a course or watches a lecture video.**
- **Batch Data (saved files):**
    - **Stored student records, course materials, past exam results, saved in OneDrive or SQL Database.**
    - **Example: Uploading a batch of last semester's grades.**

Step 2: Data Ingestion
How the data is collected:

- Stream data goes into:
    - **Azure Monitor** → Checks if any app issues happen (like a student facing login problems).
    - **Azure Log Analytics** → Records all student activities (like course enrollments).
    - **Azure Event Hubs** → Collects live events happening every second (like "student watched Lesson 5").
- Batch data is moved by:
    - **Azure Data Factory** → Picks up your saved student lists, course updates, exam results and brings them into the system automatically.

Step 3: Storage and Processing
How your university's data is organized:

- 🥉 **Bronze Layer (Raw Data)**:
    - Example: Raw messy student activity — some records are incomplete, some login fails, etc.
    - Stored in **Azure Data Lake Gen 2**.
- 🥈 **Silver Layer (Cleaned Data)**:
    - Example: After cleaning:
        - Remove duplicate student logins
        - Fix missing grades
        - Correct wrongly entered course names
    - Stored in **Azure Delta Lake**.
- 🥇 **Gold Layer (Final/Aggregated Data)**:
    - Example: Final polished reports:
        - How many students enrolled in "Data Science 101" this month?

        - Which video lessons are most watched?
    - Stored in **Azure Data Warehouse** for easy reporting.

Step 4: Consumption

How you use the data to improve YOUR Online University:

- **Azure Logic Apps**:
    - Example: Automatically send a reminder email to students who have not completed 50% of their course.
- **Azure Machine Learning**:
    - Example: Predict which students are at risk of dropping out and send them motivational emails.
- **Azure Data Share**:
    - Example: Share enrollment reports securely with your partner universities.
- **Power BI**:
    - Example: Build dashboards showing:
        - Total new enrollments each week
        - Most popular courses
        - Average student course completion rate
        - Regions where most students are signing up

5 PIPELINE DESIGN

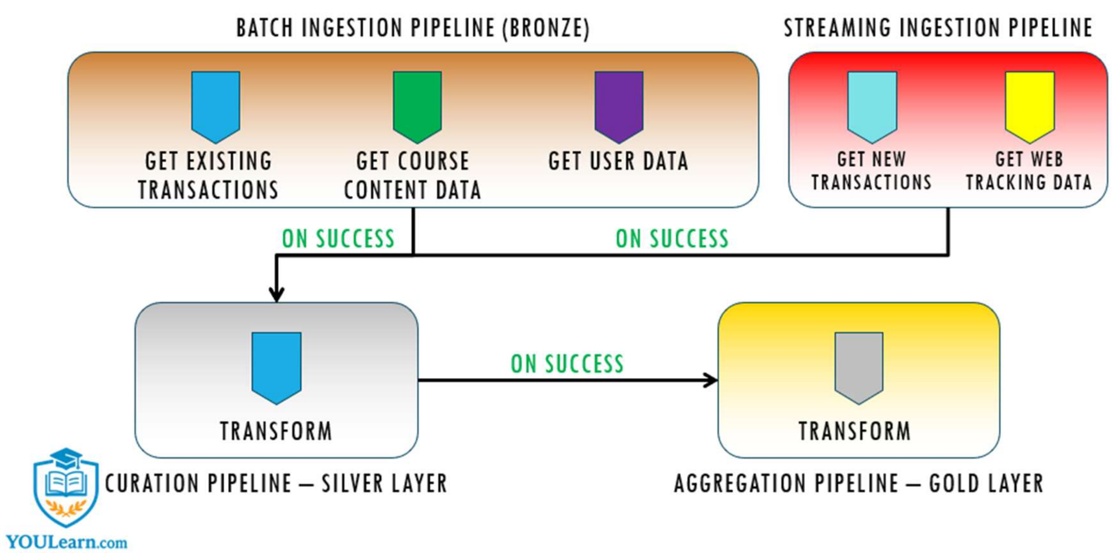

1. Real-time and Batch Data Collection

- Stream Data from **Web Apps** and **Mobile Apps** (user clicks, page visits, transactions) is captured **live** through:
    - Azure Monitor → Azure Log Analytics → Azure Event Hubs
- Batch Data like course content, user profiles, or transactional history is collected from:
    - **OneDrive**, **SQL Databases** via **Azure Data Factory**.

Outcome: we can ingest both *historical* and *real-time* data seamlessly.

2. Unified Storage and Processing

- Bronze Layer → stores *raw data* (no changes).
- Silver Layer → stores *cleaned, curated data* (after transformation).
- Gold Layer → stores *aggregated data* (ready for analytics and reporting).

Outcome: Clean separation between raw, processed, and analytics-ready data — making it easy to debug, audit, and reprocess if needed.

3. Advanced Data Processing & Transformation

- Using **Azure Synapse Analytics** and **Data Factory**, you automate:
    - Data cleaning
    - Joining different datasets (e.g., user behavior + course content)
    - Aggregation (e.g., course popularity, student success rates)

Outcome: Faster ETL (Extract, Transform, Load) pipelines, with scalable processing power.

4. Machine Learning and Advanced Analytics Integration

- Curated and Aggregated data can be pushed to:
    - **Azure Machine Learning** to build AI models (e.g., predict course completion, recommend courses).
    - **Azure Data Share** to collaborate data securely across teams or departments.
    - **Azure Logic Apps** to automate business workflows (e.g., notify users about new courses based on browsing history).

Outcome: Smarter predictions, personalization, and automation based on clean, up-to-date data.

5. Business Intelligence and Reporting

- Gold layer feeds into **Power BI** dashboards.

- Real-time KPIs (e.g., number of new students, revenue per course).
- Deep analytics (e.g., churn rate analysis, top-performing courses).

Outcome: Decision-makers can make fast, data-driven decisions with beautiful, interactive dashboards.

6 Pipeline approach

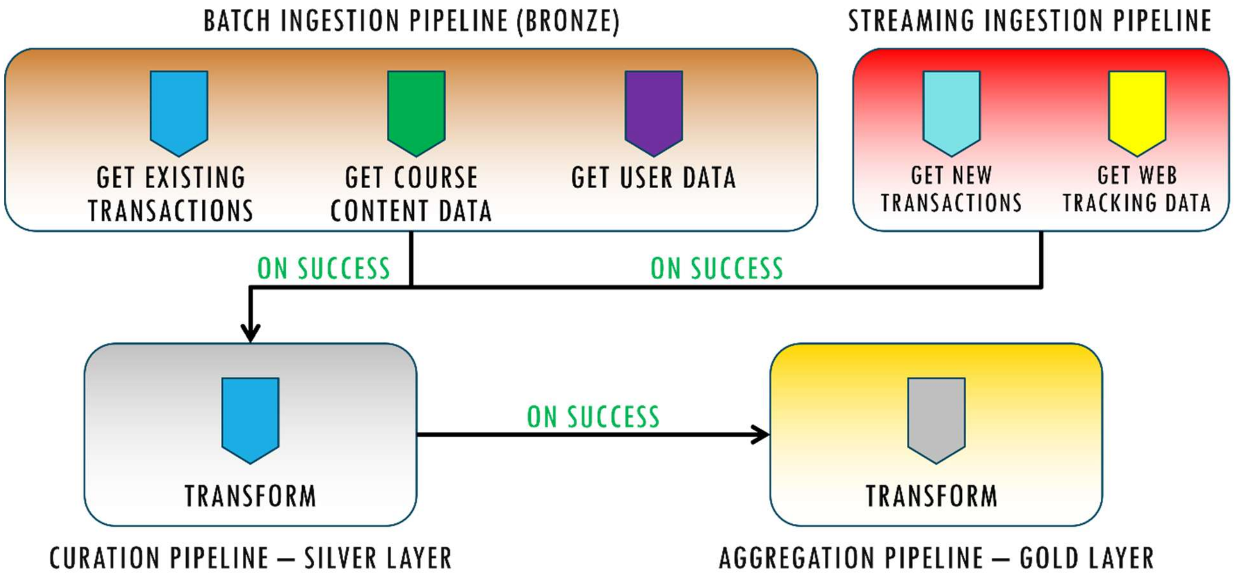

Event-based approach

- **Master Pipeline**: A controller pipeline designed in **Azure Data Factory** to monitor and trigger sub-pipelines.

  Sub-Pipelines: Handle batch ingestion, curation, and aggregation tasks.

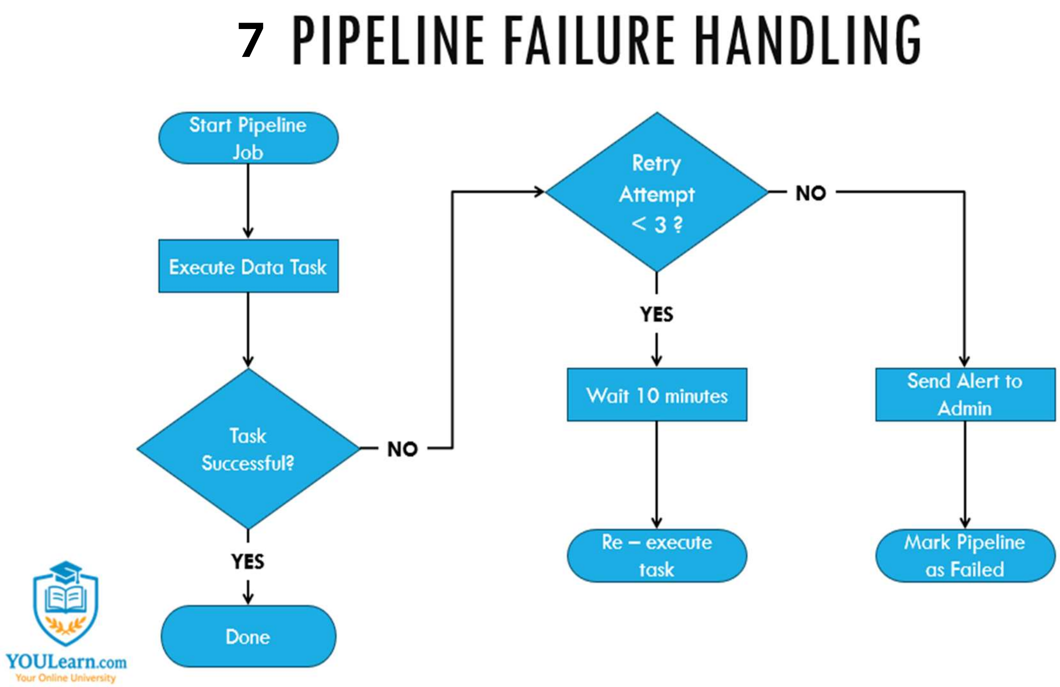

Pipeline Failure Handling

We implemented robust failure handling mechanisms to make pipelines fault-tolerant and self-healing:

🔁 **Retry Strategy:**

- **Retry Attempts:** Each pipeline is configured with **up to 3 retry attempts** upon failure.
- **Retry Interval:** A delay of **10 minutes** between retries to allow for transient issues (e.g., service outages, slow blob upload) to resolve.

If retries still fail, the master pipeline:

- Sends an alert using **Azure Monitor** and **Logic Apps**
- Triggers **incident workflows** (e.g., notifying engineering teams, logging error into tracking systems)

This ensures high pipeline uptime and reduces operational bottlenecks.

8 CONCLUSION:

- **Smarter architecture**: Replaces manual processes with an automated, scalable data pipeline.
- **Layered design**: Bronze–Silver–Gold model ensures clean, curated, and actionable data.
- **Azure-powered**: Leverages Data Factory, Event Hub, Synapse, and Power BI for full data lifecycle.
- **Real-time insights**: Enables faster decisions and personalized learning through ML.
- **Resilient system**: Event-driven pipeline with retries and proactive monitoring ensures reliability.
- **Future-ready**: Scalable foundation for continued growth and innovation.